



Fondamenti di Analisi dei Dati

from **data analysis** to **predictive techniques**

Prof. Antonino Furnari (antonino.furnari@unict.it)

Corso di Studi in Informatica

Dip. di Matematica e Informatica

Università di Catania



Università
di Catania

Beyond Linear Regression

Extending the linear model to handle complex, non-linear relationships while
managing the critical trade-off between bias and variance in predictive modeling.

The Journey Ahead

01

Statistical Fixes

Extending the linear model with interaction terms and polynomial features

03

Machine Learning Perspective

Shifting focus from understanding to prediction accuracy

05

Regularization Solutions

Ridge and Lasso techniques to control model complexity

02

Quadratic & Polynomial Regression

Capturing non-linear patterns in data through higher-degree terms

04

Overfitting Problem

Understanding when models become too powerful

06

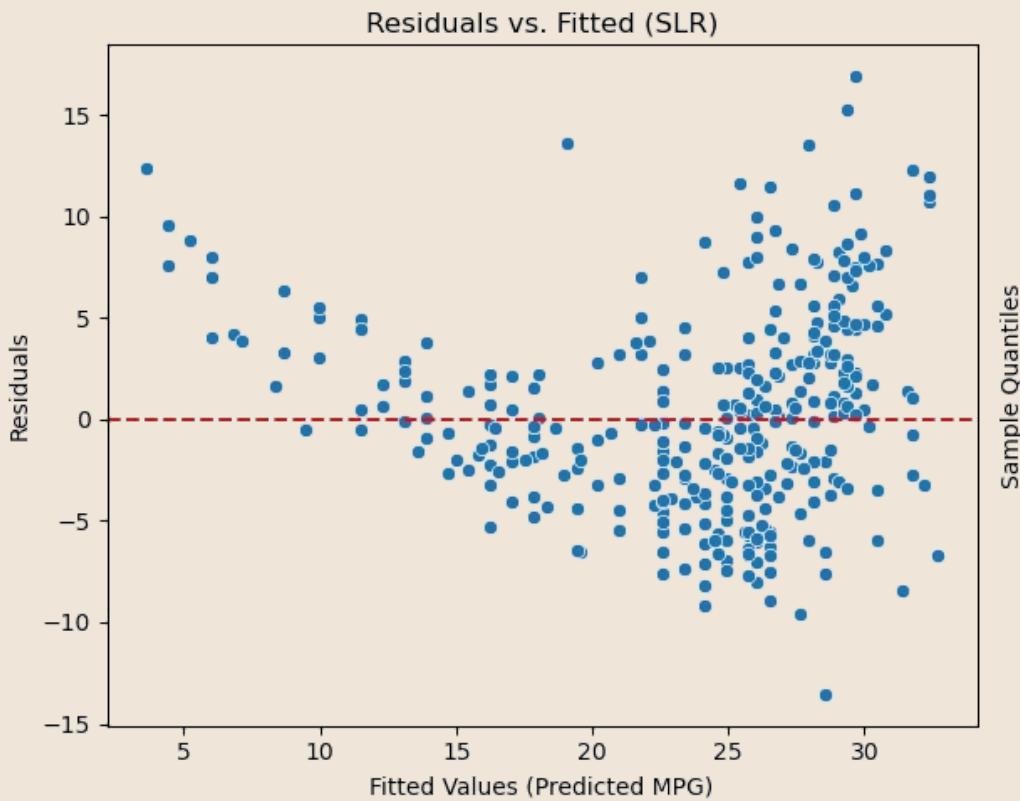
Practical Implementation

The scikit-learn workflow for real-world applications

Problems with Simple Linear Models

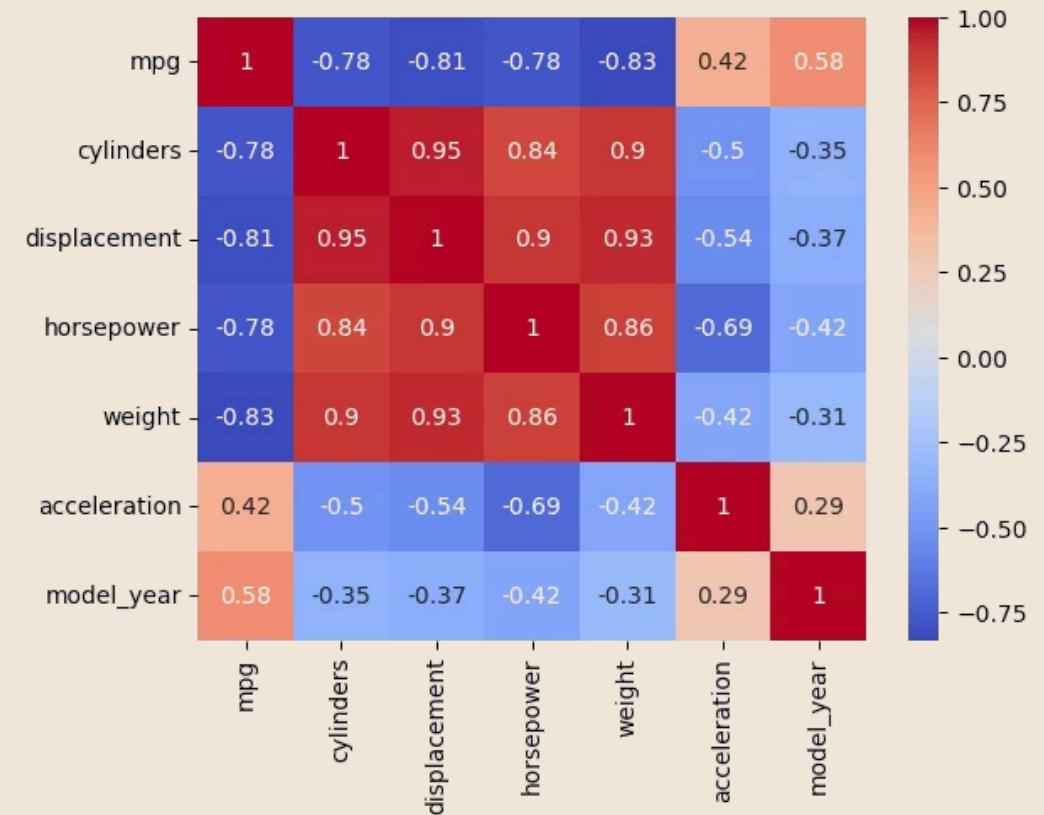
1. Non-Linear Patterns

When modeling $\text{mpg} \sim \text{horsepower}$, residual plots revealed a clear U-shaped pattern. This indicates **underfitting** (high bias) because the real-world relationship isn't linear.



2. Multicollinearity

Using multiple correlated predictors like horsepower and weight caused unstable p-values and confused interpretation. High correlation between predictors makes coefficient estimates unreliable.



Today we'll fix these problems by extending the linear model, though we'll sacrifice some interpretability in exchange for predictive power.

Interaction Terms: When Effects Depend on Context

Simple linear models assume additive effects. But what if the effect of horsepower on MPG depends on the car's weight?

Additive Model

$$mpg = \beta_0 + \beta_1 \text{horsepower} + \beta_2 \text{weight}$$

Effect of horsepower is constant: β_1

Adj. R-squared: 0.7049

	coef	std err	t	P> t	[0.025	0.975]
Intercept	39.9359	0.717	55.660	0.000	38.525	41.347
horsepower	-0.1578	0.006	-24.489	0.000	-0.171	-0.145

Interaction Model

$$mpg = \beta_0 + \beta_1 \text{horsepower} + \beta_2 \text{weight} + \beta_3 (\text{horsepower} \times \text{weight})$$

Effect varies: $\beta_1 + \beta_3 \times \text{weight}$

Adj. R-squared: 0.7465 ✓

	coef	std err	t	P> t	[0.025	0.975]
Intercept	63.5579	2.343	27.127	0.000	58.951	68.164
horsepower	-0.2508	0.027	-9.195	0.000	-0.304	-0.197
weight	-0.0108	0.001	-13.921	0.000	-0.012	-0.009
horsepower:weight	5.355e-05	6.65e-06	8.054	0.000	4.05e-05	6.66e-05

Interpretation of the Interaction Model

How to interpret the model? We can rewrite it as:

$$mpg = \beta_0 + (\beta_1 + \beta_3 \times \text{weight}) \times \text{horsepower} + \beta_2 \text{weight}$$

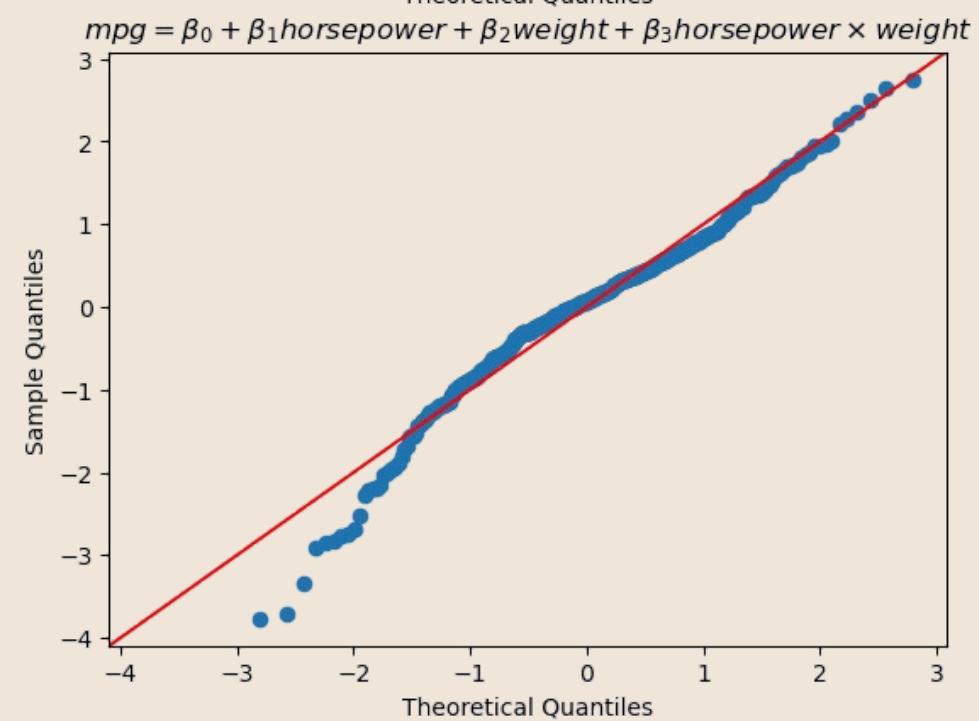
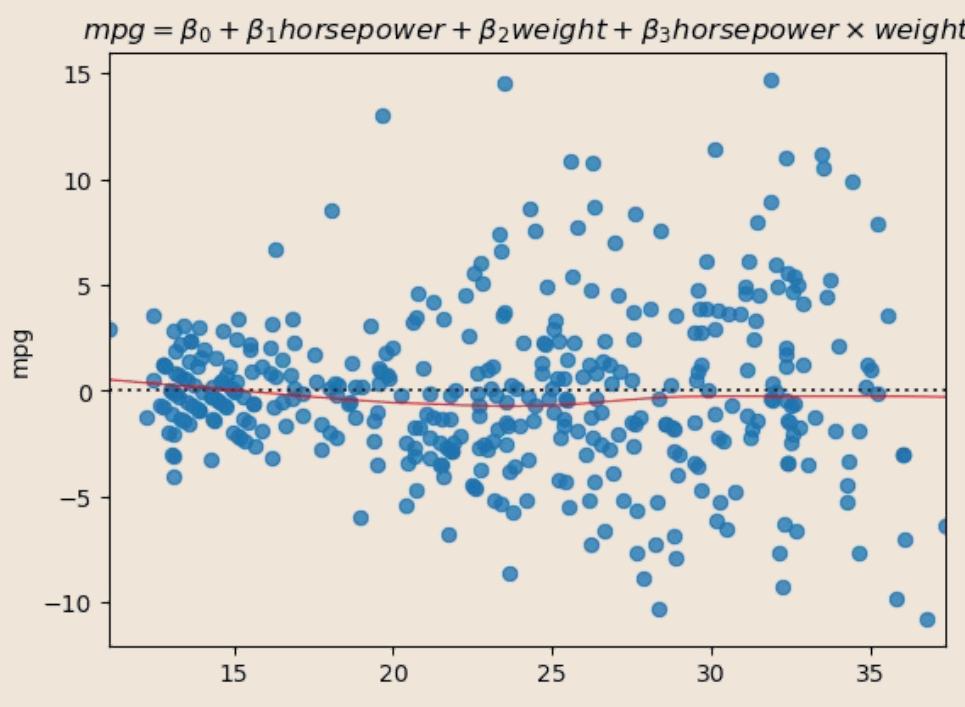
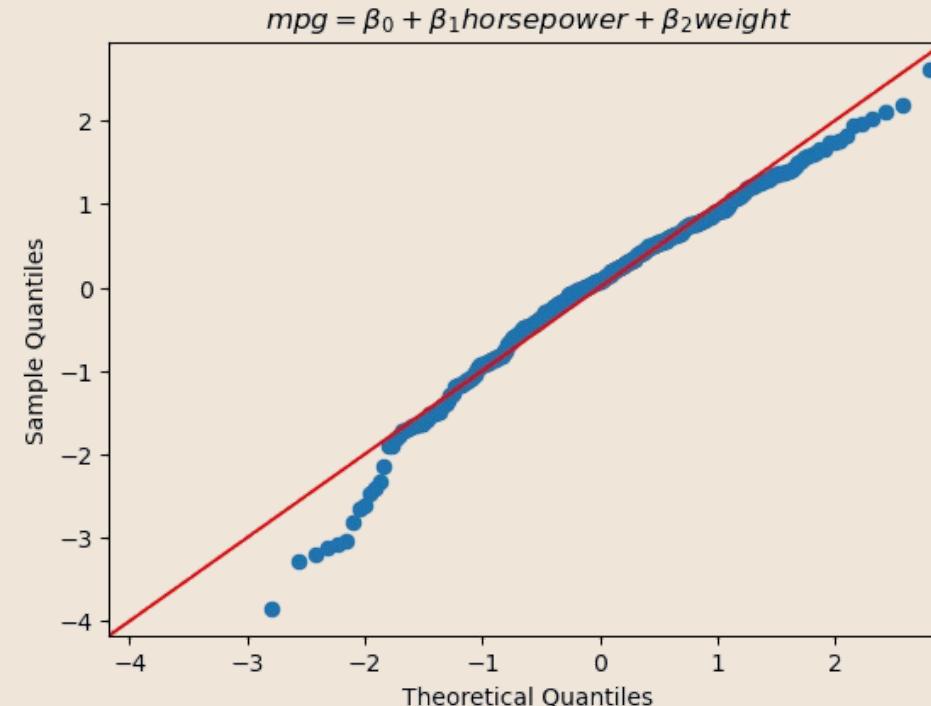
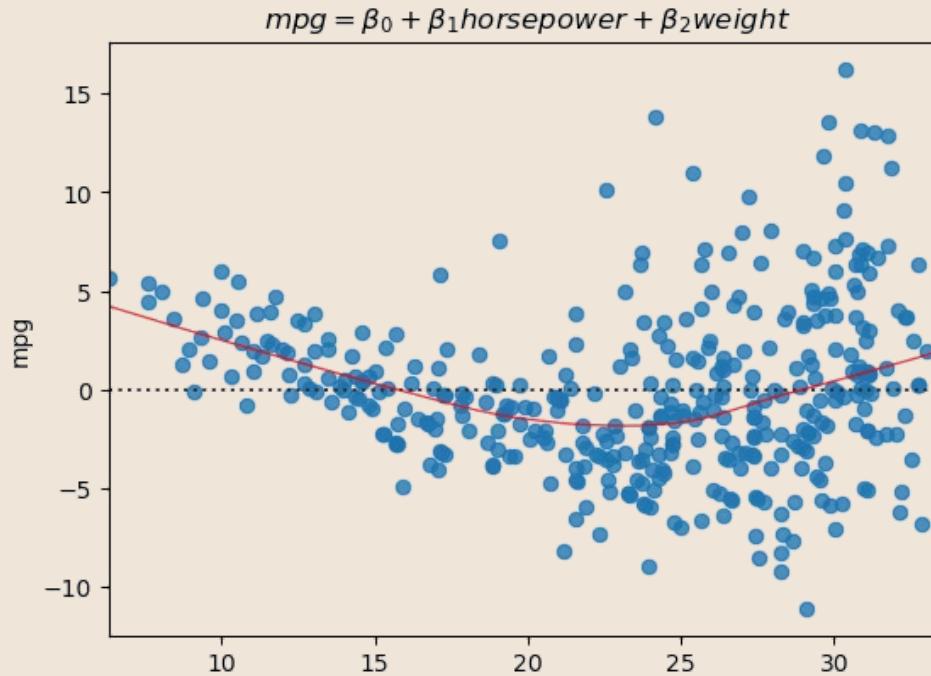
The effect of horsepower is no longer a single number, β_1 . It is now $(\beta_1 + \beta_3 \times \text{weight})$. This means that the effect of horsepower **depends on** the weight of the car.

In practice, **when we increase weight by one unit, we increase the effect of horsepower on mpg by β_3 units**. This means that for heavier cars, horsepower has a less negative effect on mpg.

This makes sense mechanically: in a very heavy car, horsepower is needed just to move it, so the efficiency penalty is smaller compared to a light car where extra horsepower is “overkill.”

Comparing the new Model to a Simple Regressor

The model is a little better!

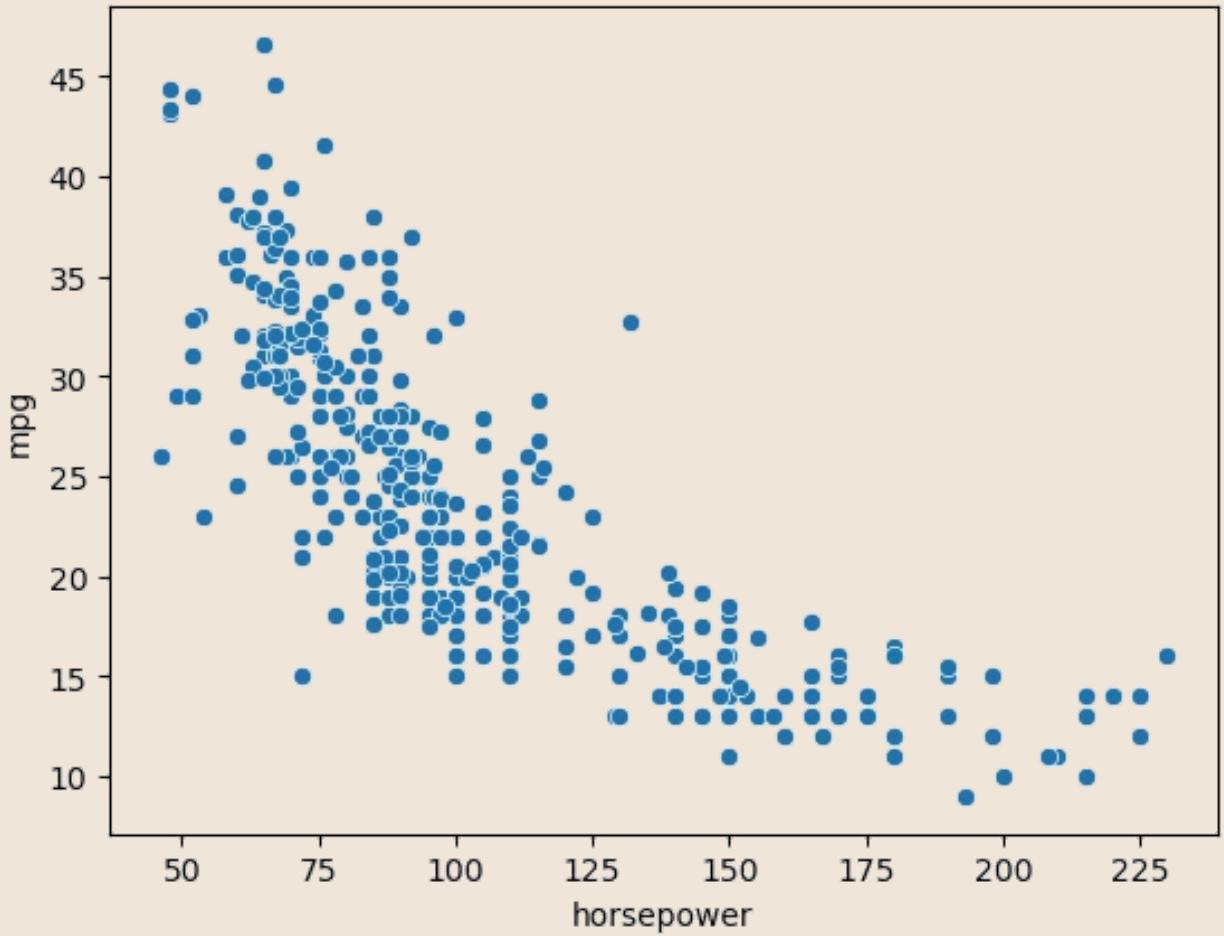


The Non-Linear Reality

The Problem

The relationship between MPG and horsepower isn't linear—it's curved. A straight line simply can't capture this pattern.

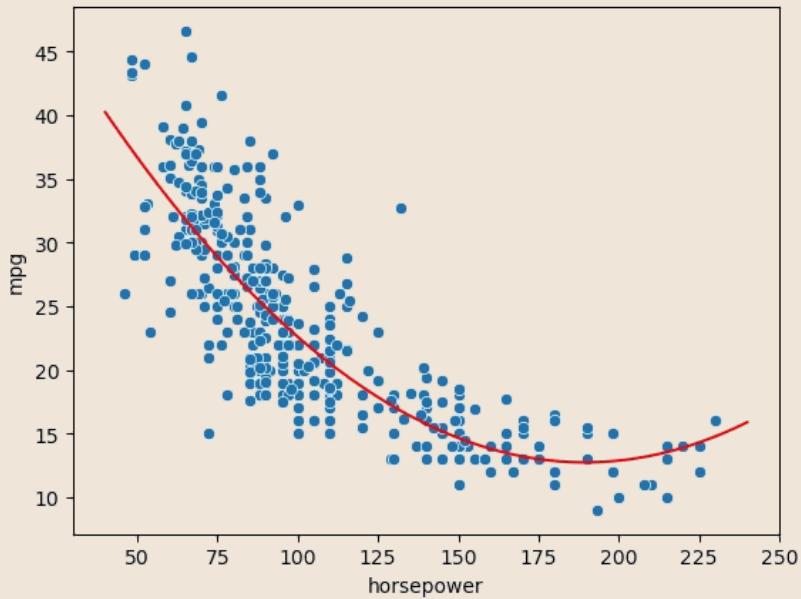
We need a model that can bend and curve to fit the data's natural shape.



Quadratic Regression: Adding Curvature

We can capture the curved by adding a squared term:

$$mpg = \beta_0 + \beta_1 \text{horsepower} + \beta_2 \text{horsepower}^2$$



	coef	std err	t	P> t	[0.025	0.975]
Intercept	56.9001	1.800	31.604	0.000	53.360	60.440
horsepower	-0.4662	0.031	-14.978	0.000	-0.527	-0.405
I(horsepower ** 2)	0.0012	0.000	10.080	0.000	0.001	0.001

This model has a higher R^2 of 0.688, larger than the base 0.608.

We now have the model:

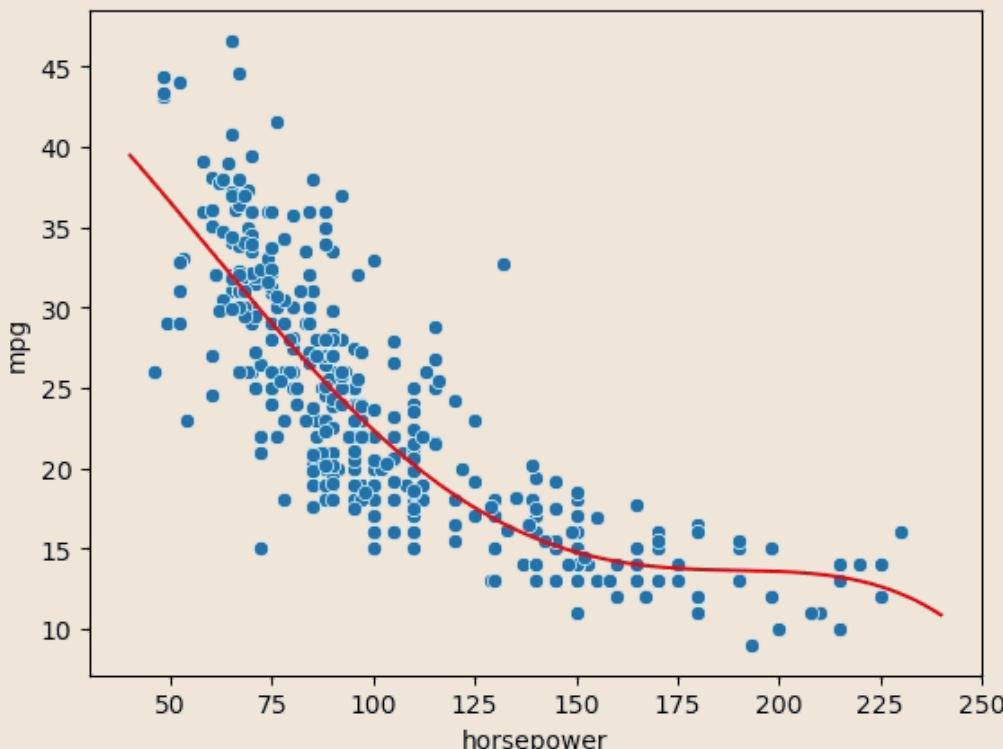
$$y = 58.9001 - 0.4662x + 0.0012x^2$$

- ❑ While we have a better model, we lost interpretability. How do we interpret the coefficient of the quadratic term? How does the interpretation of other coefficients change?

Polynomial Regression: More Flexibility

We can extend this idea to higher-degree polynomials. For degree $d=4$:

$$mpg = \beta_0 + \beta_1x + \beta_2x^2 + \beta_3x^3 + \beta_4x^4$$



Key Insight

While the polynomial function is non-linear in the variable, it's still **linear in its coefficients**. This means we can solve it using ordinary least squares!

We simply create new variables for x^2 , x^3 , etc., and treat it as a linear regression problem.

For instance, this model:

$$y = \beta_0 + \beta_1x + \beta_2y$$

becomes in the case of degree $d = 2$:

$$y = \beta_0 + \beta_1x + \beta_2y + \beta_3x^2 + \beta_4y^2 + \beta_5xy$$

The Interpretability Trade-Off

What is the "effect" of horsepower on MPG?

Linear Model

Simple answer: β_1

One number tells the whole story. Easy to explain and understand.

Quadratic Model

The effect *depends on the value of horsepower*. No single number captures it.

This is the pivot point. If we're willing to sacrifice simple interpretability for better fit, why stop here? This is where we cross from **statistical inference** into **Machine Learning**.

Shifting Perspectives: From Understanding to Prediction



Statistical Approach

- Focus on p-values and R²
- Interpret coefficients
- Understand relationships
- Use training data metrics



Machine Learning Approach

- Focus on predictive accuracy
- Optimize for unseen data
- Embrace complexity
- Use test set metrics

When prediction is the goal, we only care about performance on **new, unseen data** (the test set). This requires new metrics and a new mindset.

Metrics for Prediction

1

Mean Squared Error (MSE)

$$MSE_{test} = \frac{1}{m} \sum_{j=1}^m (y_j - \hat{y}_j)^2$$

Average of squared errors. Problem: units are squared (e.g., meters²), making interpretation difficult.

2

Root Mean Squared Error (RMSE)

$$RMSE_{test} = \sqrt{MSE_{test}}$$

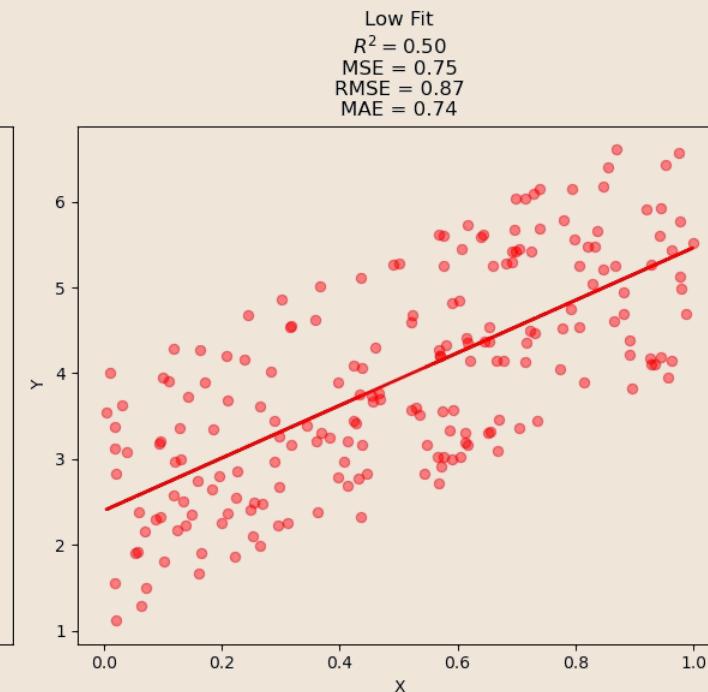
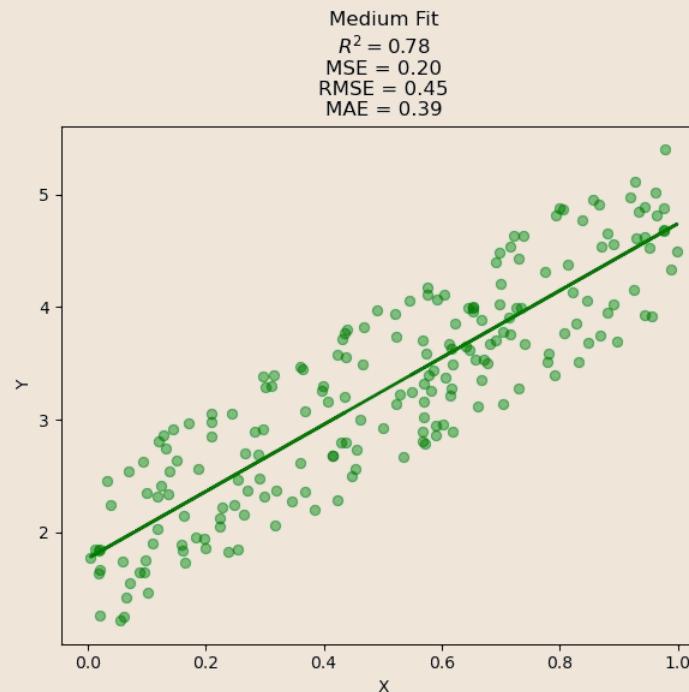
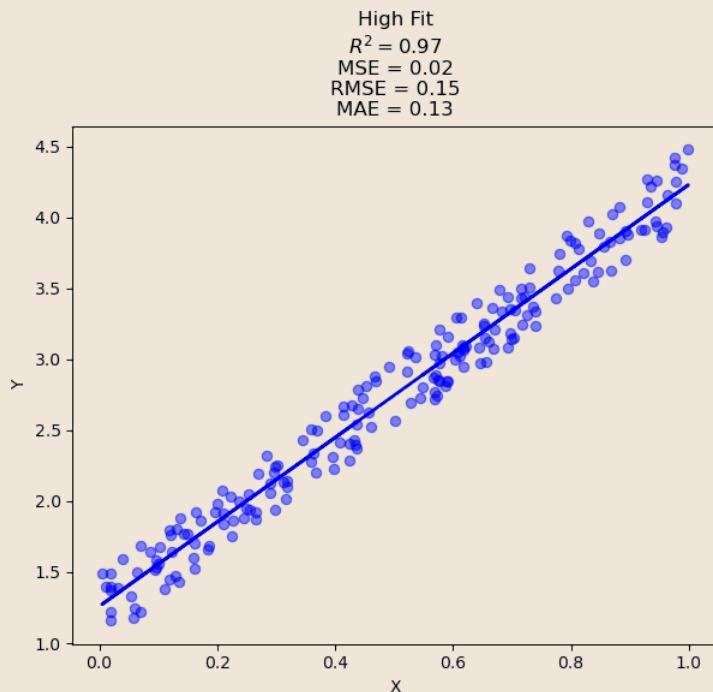
Most common metric. Same units as Y.
Penalizes large errors more heavily.
Interpretation: "On average, predictions are wrong by [RMSE value]."

3

Mean Absolute Error (MAE)

$$MAE_{test} = \frac{1}{m} \sum_{j=1}^m |y_j - \hat{y}_j|$$

Average absolute error. Same units as Y.
More robust to outliers than RMSE—doesn't disproportionately penalize large errors.



The Overfitting Danger

If degree=2 is good, is degree=20 better?

Underfitting

Linear model is too simple. Misses the pattern.

High bias.

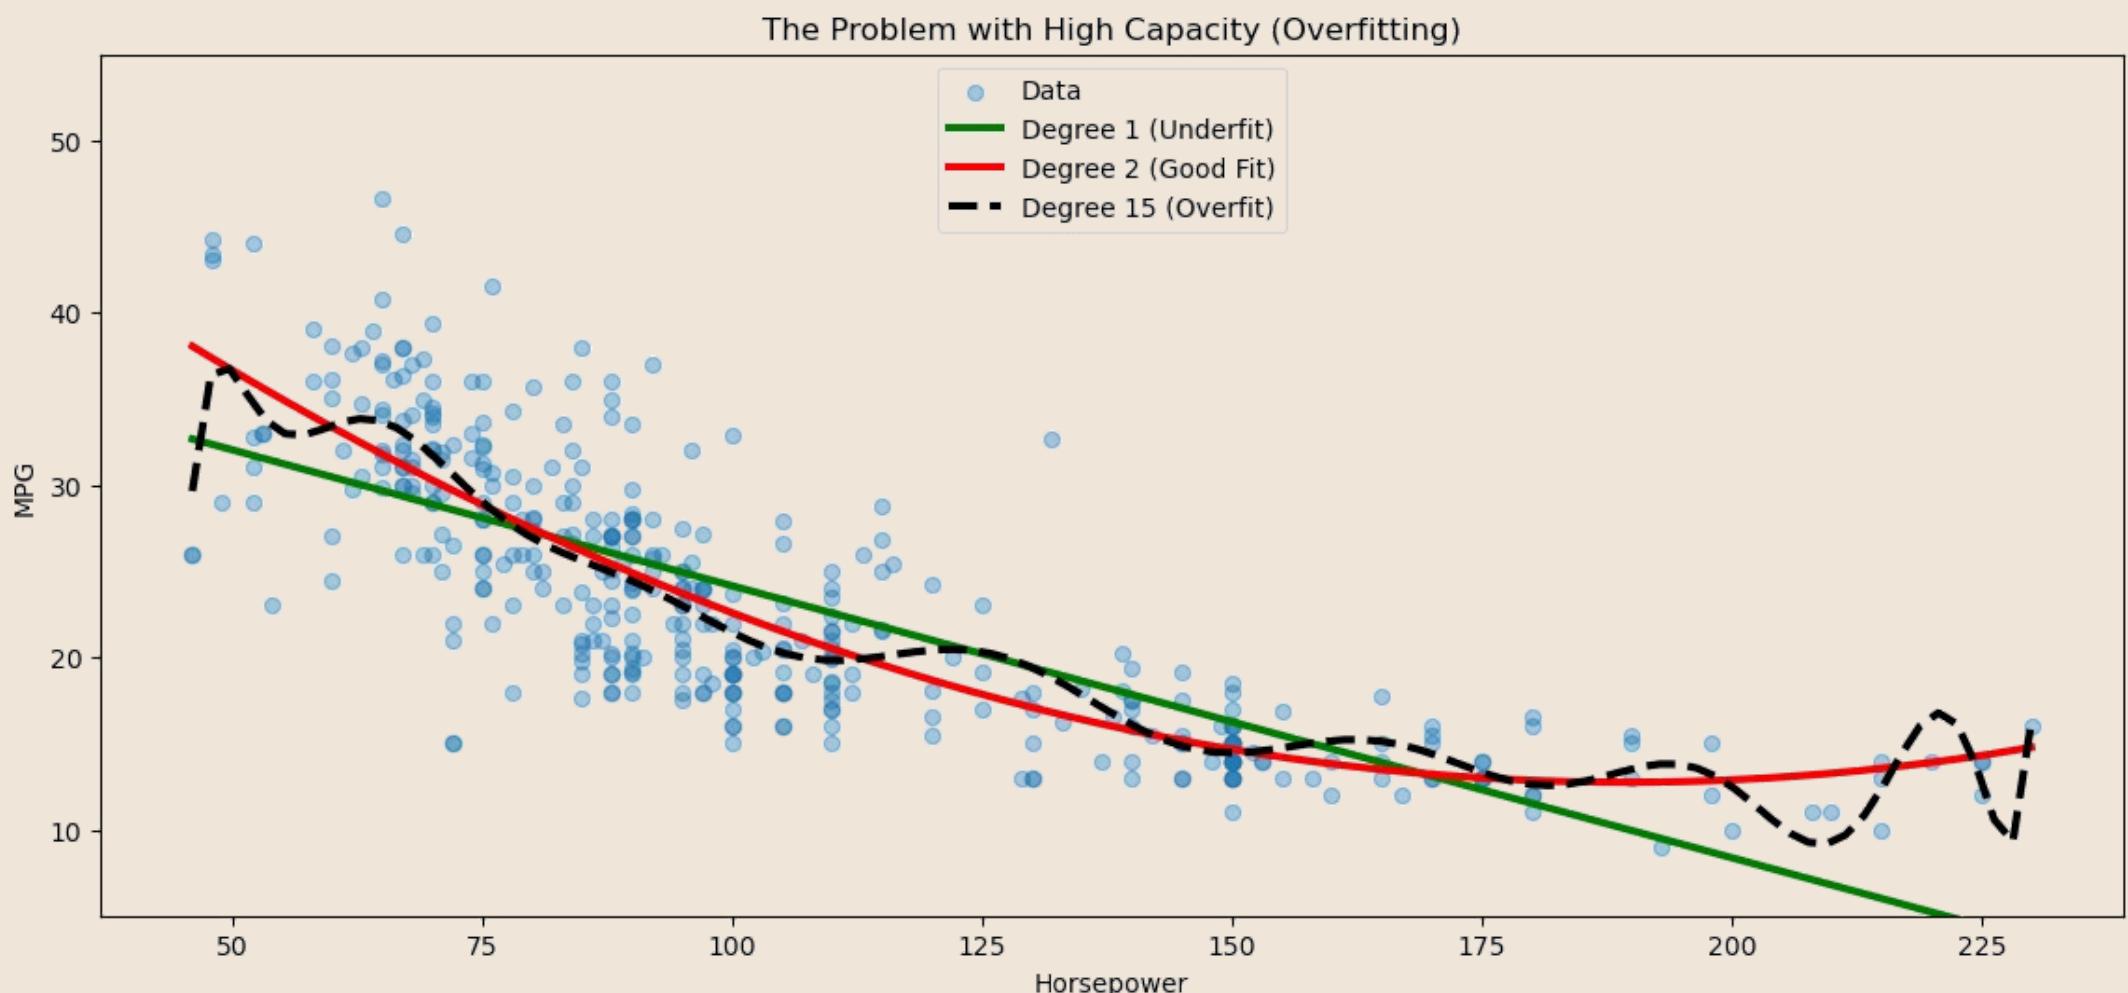
Just Right

Quadratic model captures the signal. Balanced bias-variance.

Overfitting

High-degree polynomial learns the noise.

High variance.



- A high-degree polynomial can "wiggle" to fit every single training point, achieving perfect R^2 on training data but failing miserably on new data.

The Solution: Regularization

Regularization adds a penalty to the cost function to discourage model complexity:

$$RSS + \text{Penalty Term}$$



A "wiggly" overfit curve requires massive β coefficients. By penalizing large coefficients, we force the model to find a simpler, smoother curve.

- ❑ **Bonus:** Regularization also solves multicollinearity by stabilizing the coefficient estimates!

Ridge Regression (L_2 Penalty)

Ridge adds a penalty proportional to the **sum of squares** of coefficients:

$$RSS_{L2} = \sum_{i=1}^m (y_i - \hat{y}_i)^2 + \lambda \sum_{j=1}^n \beta_j^2$$

This is solved with OLS as follows:

$$\hat{\beta}_{Ridge} = (X^T X + \lambda I)^{-1} X^T y$$

where the λI term improves conditioning (invertibility) of the matrix even in the presence of multicollinearity.

$\lambda = 0$

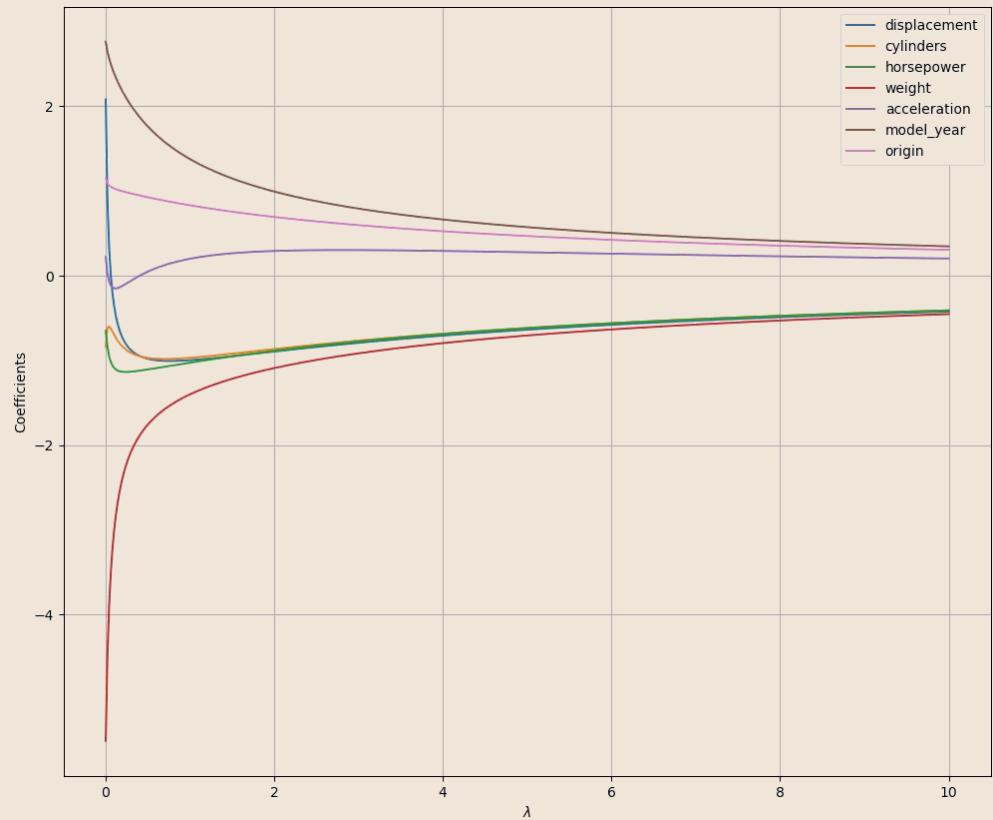
Standard OLS. No penalty.
Risk of overfitting.

$\lambda = \text{optimal}$

Sweet spot. Balanced bias-variance trade-off.

$\lambda \rightarrow \infty$

All coefficients forced to zero. Flat line. Underfitting.



- **Critical:** Standardize (z-score) all features first! Otherwise, features with different scales will be penalized unfairly.

Interpreting Ridge Coefficients

Variable	Ridge	OLS
displacement	-0.998	2.079
cylinders	-0.970	-0.841
horsepower	-1.027	-0.652
weight	-1.402	-5.492
acceleration	0.200	0.222
model_year	1.376	2.762
origin	0.830	1.147

- ❑ Ridge parameters have smaller magnitude due to regularization. They can't be interpreted statistically like OLS coefficients. Instead, they indicate **relative importance** of variables for prediction.

Lasso Regression (L_1 Penalty)

Lasso uses the **sum of absolute values** of coefficients:

$$RSS_{L1} = \sum_{i=1}^m (y_i - \hat{y}_i)^2 + \lambda \sum_{j=1}^n |\beta_j|$$

Ridge vs. Lasso

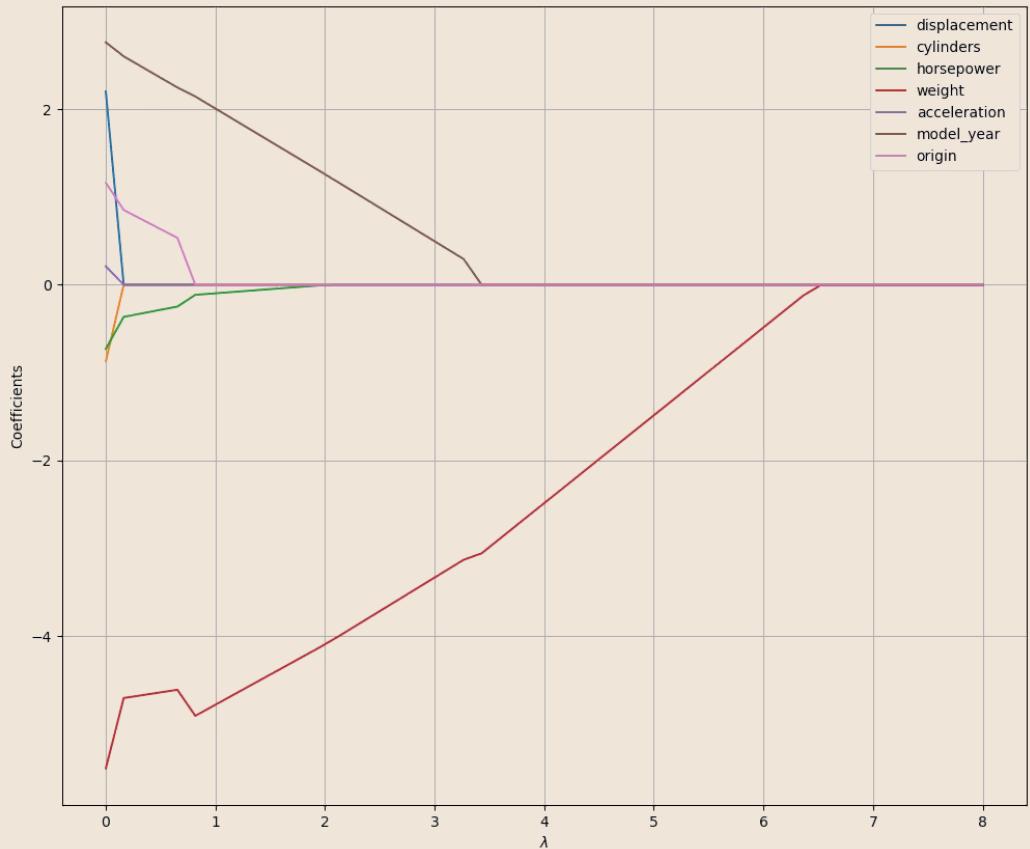
- **Ridge:** Shrinks coefficients smoothly
- **Lasso:** Sets coefficients exactly to zero

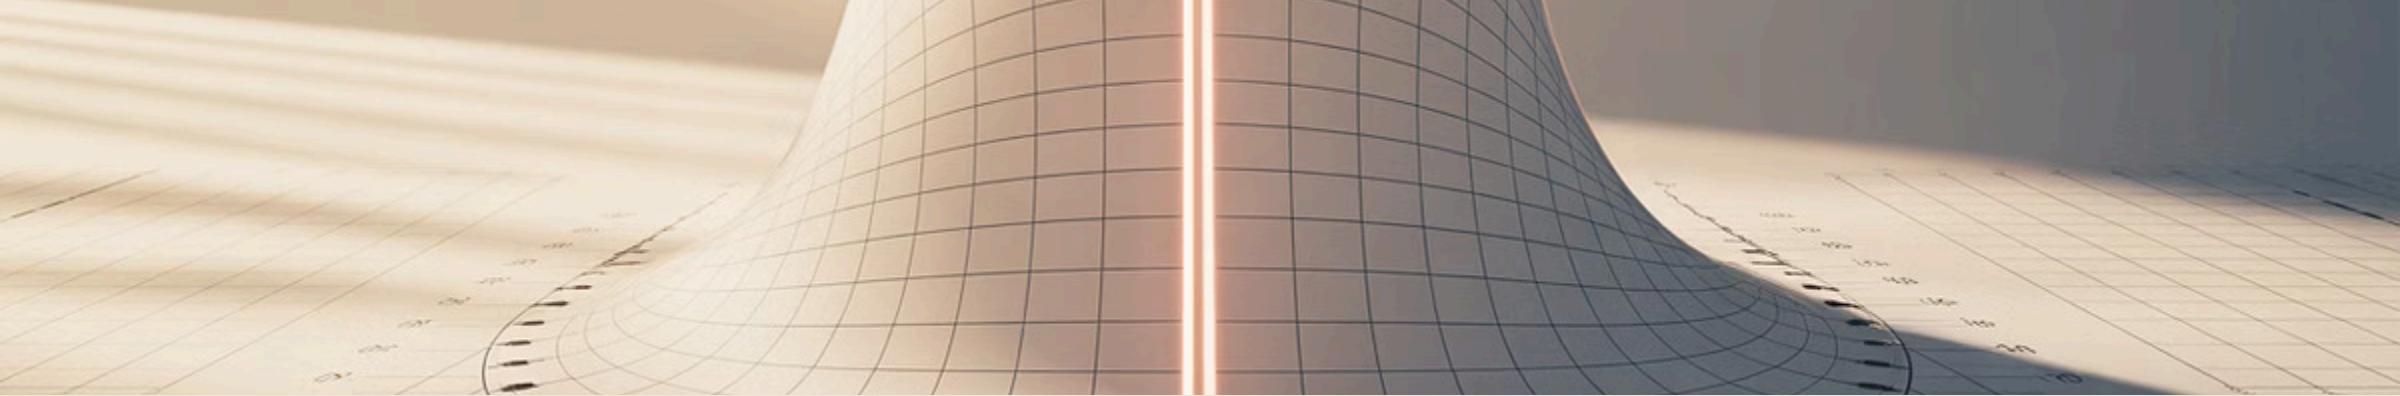
The "sharp corners" of the absolute value function allow Lasso to eliminate features entirely.

Lasso's Characteristic

Automatic feature selection!

When you have hundreds of features, **Lasso identifies which ones matter and removes the rest**. This creates simpler, more interpretable models.





The Bias-Variance Trade-Off

$\lambda = 0$ (No Penalty)

Low Bias, High Variance

Overfitting: perfect fit on training data, poor on test data

$\lambda \rightarrow \infty$ (Heavy Penalty)

High Bias, Low Variance

Underfitting: flat line, poor fit everywhere

1

2

3

$\lambda = \text{Optimal}$

Balanced

Sweet spot: lowest total error on test data

λ is our "tradeoff knob"—the entire goal of machine learning is finding the optimal λ that minimizes test error.

Validation Curve: Finding the Sweet Spot

Left Side ($\lambda \approx 0.001$)

Training error very low, test error high.
Overfitting (high variance region).

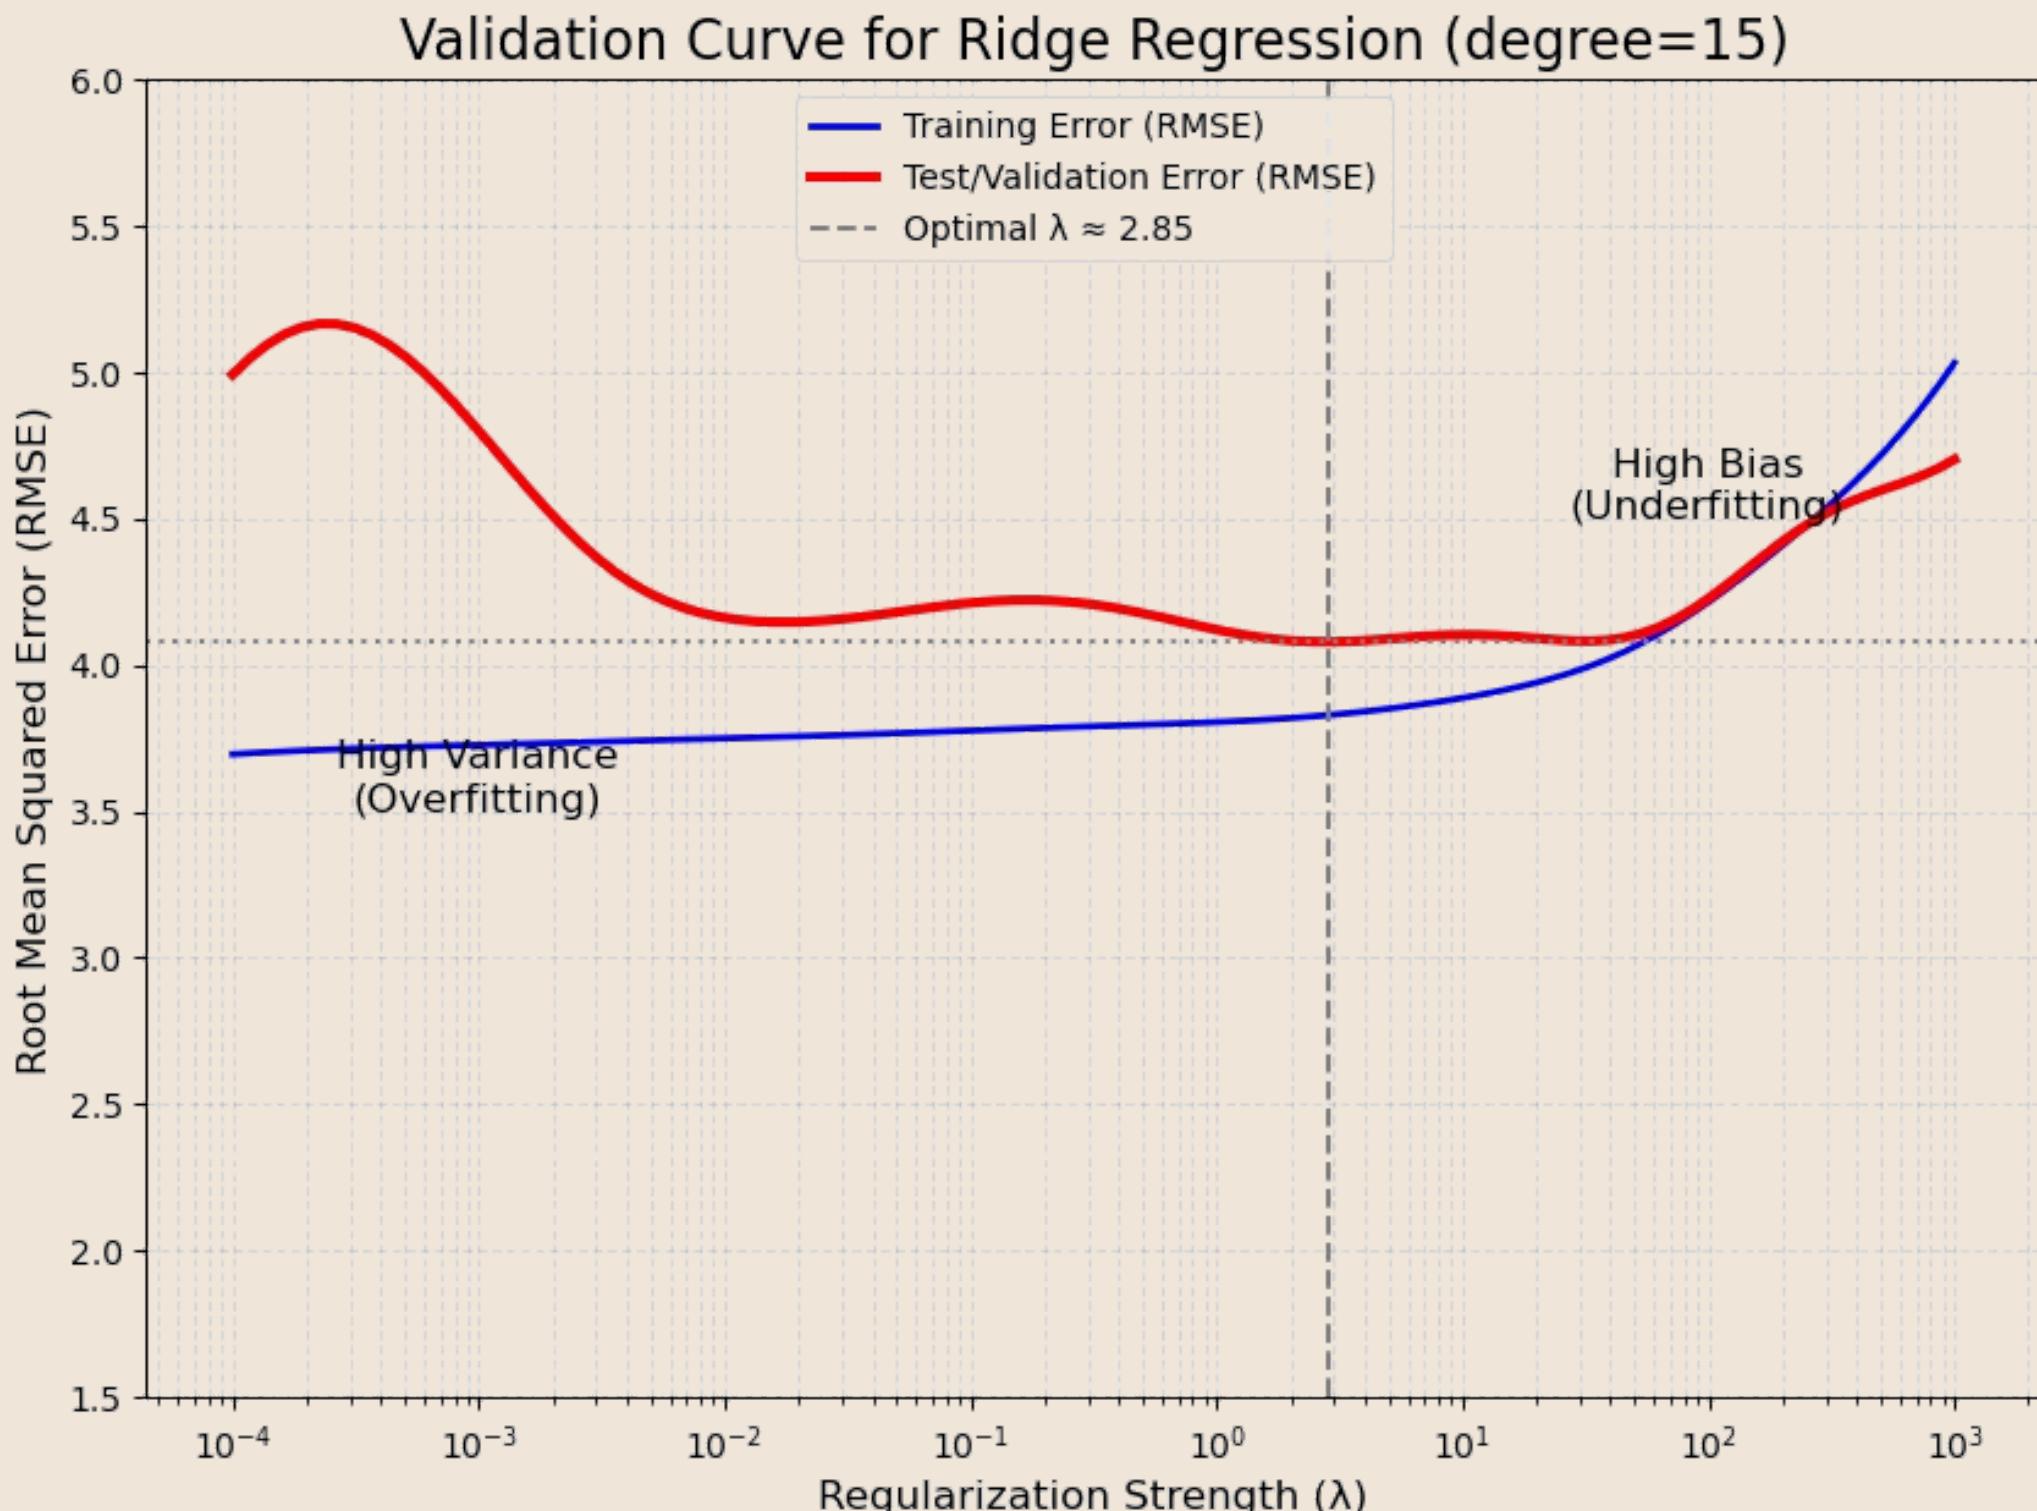
Sweet Spot (Dashed Line)

Test error at minimum. **Optimal λ** —perfect balance between bias and variance.

Right Side ($\lambda > 1000$)

Both errors high and similar. **Underfitting** (high bias region).

This visualization shows how regularization controls the bias-variance trade-off for a degree-15 polynomial. Cross-validation helps us find the optimal λ systematically.

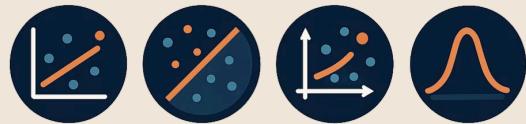


Linear Regression in Python



	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin
0	18.0	8	307.0	130.0	3504	12.0	70	usa
1	15.0	8	350.0	165.0	3693	11.5	70	usa
2	18.0	8	318.0	150.0	3436	11.0	70	usa
3	16.0	8	304.0	150.0	3433	12.0	70	usa
4	17.0	8	302.0	140.0	3449	10.5	70	usa

Conclusions and Next Steps



We Have Explored:

- Interaction Terms
- Quadratic Model
- Polynomial Model
- Regularization

In the next lectures, we will look at classification problems

References

- Chapter 3 of [1]
- Parts of chapter 11 of [2]

[1] Heumann, Christian, and Michael Schomaker Shalabh. Introduction to statistics and data analysis. Springer International Publishing Switzerland, 2016.

[2] James, Gareth Gareth Michael. An introduction to statistical learning: with applications in Python, 2023. <https://www.statlearning.com>