



# Fondamenti di Analisi dei Dati

from **data analysis** to **predictive techniques**

Prof. Antonino Furnari ([antonino.furnari@unict.it](mailto:antonino.furnari@unict.it))

Corso di Studi in Informatica

Dip. di Matematica e Informatica

Università di Catania




Università  
di Catania

## Multiclass Logistic Regression

Extending logistic regression beyond binary classification to handle multiple classes, and exploring the algorithm from a purely predictive machine learning perspective.

# The Diabetes Dataset

In many cases, we may need to study the relationship between some dependent variables and an independent one which is **categorical with more than two possible levels**. Let's consider the diabetes dataset:



# Multinomial Logistic Regression

The multinomial logistic regression model extends binary logistic regression by modeling the log-odds ratio between each class k and a baseline class (typically class 1):

$$\log \left( \frac{P(Y = k|X = \mathbf{x})}{P(Y = 1|X = \mathbf{x})} \right) = \beta_{k0} + \beta_{k1}x_1 + \dots + \beta_{kn}x_n$$

The model estimates  $(n+1) \times (K-1)$  parameters, where n is the number of predictors and K is the number of classes. Each non-baseline class gets its own set of coefficients.

Through mathematical derivation, we can express the class probabilities as:

## Non-baseline classes

$$P(Y = k|X = \mathbf{x}) = \frac{e^{\beta_k^T \mathbf{x}}}{1 + \sum_{l=2}^K e^{\beta_l^T \mathbf{x}}}$$

## Baseline class

$$P(Y = 1|X = \mathbf{x}) = \frac{1}{1 + \sum_{l=2}^K e^{\beta_l^T \mathbf{x}}}$$

These expressions compute class probabilities once parameters are estimated and predictors are observed.

- **Optimisation** We won't see the details, but this is optimised by establishing a similar objective (maximum likelihood or cross-entropy).

# Multinomial Logistic Regression Example

If we fit this model with statsmodels:

$$\text{group} = \beta_0 + \beta_1 \text{relwt} + \beta_2 \text{glutest}$$

We obtain the following output:

Optimization terminated successfully.  
Current function value: 0.100843  
Iterations 14

## MNLogit Regression Results

Dep. Variable:	group	No. Observations:	145
Model:	MNLogit	Df Residuals:	141
Method:	MLE	Df Model:	2
Date:	Sat, 15 Nov 2025	Pseudo R-squ.:	0.9013
Time:	15:53:12	Log-Likelihood:	-14.622
converged:	True	LL-Null:	-148.10
Covariance Type:	nonrobust	LLR p-value:	1.076e-58

group=1	coef	std err	z	P> z	[0.025	0.975]
Intercept	-90.3738	42.527	-2.125	0.034	-173.726	-7.022
glutest	0.2153	0.101	2.128	0.033	0.017	0.413

group=2	coef	std err	z	P> z	[0.025	0.975]
Intercept	-109.8266	43.021	-2.553	0.011	-194.147	-25.506
glutest	0.2471	0.102	2.428	0.015	0.048	0.447

Model Likelihood Ratio Test:  
266.9538631863749 1.0757346252010753e-58

## Model Overview


<b>Model Convergence</b> Optimization terminated successfully after 14 iterations with log-likelihood of -14.622	<b>Likelihood Ratio Test</b> Statistic: 266.95 with p-value ≈ 0, showing highly significant improvement over null model	<b>Pseudo R<sup>2</sup></b> This is $R^2 = 1 - \frac{-14.622}{-148.10} = 0.9013$ . The multinomial logistic regressor explains about 90% of the uncertainty in the outcomes.	<b>Two Coefficient Sets</b> Separate coefficients estimated for Chemical_Diabetic and Overt_Diabetic groups versus Normal baseline
---	--	---	---

## Coefficient Interpretation



<b>Chemical Diabetic Intercept: -90.37</b> When glutest is zero, odds of being Chemical_Diabetic vs Normal are $e^{-90.37} \approx 0$ , indicating extremely low probability. The subject is almost certainly Normal at baseline.	<b>Overt Diabetic Intercept: -109.83</b> When glutest is zero, odds of being Overt_Diabetic vs Normal are $e^{-109.83} \approx 0$ , an even lower probability than Chemical_Diabetic.	<b>Chemical Diabetic Slope: 0.2153</b> Each one-unit increase in glutest multiplies the odds by $e^{0.2153} \approx 1.24$ , representing a 24% increase in odds of Chemical_Diabetic vs Normal.	<b>Overt Diabetic Slope: 0.2471</b> Each one-unit increase in glutest multiplies the odds by $e^{0.2471} \approx 1.28$ , representing a 28% increase in odds of Overt_Diabetic vs Normal.
--	--	--	--

# Geometrical Interpretation of the Coefficients of a Logistic Regressor



While the logistic function itself is non-linear, **logistic regression is fundamentally a linear classifier**. The decision boundary—where  $P(y=1|x) = 0.5$ —forms a hyperplane in feature space.

The figure on the left shows an example. A linear classifier would find a line separating the two sets of points.

We will see that the logistic regressor does just that.

We define the decision boundary as the set of points of maximum uncertainty, i.e.,  $P(y = 1|\mathbf{x}) = 0.5$ . Recalling that in the case of two dimensions:

$$P(y = 1|\mathbf{x}) = \frac{1}{1 + e^{\beta_0 + \beta_1 x_1 + \beta_2 x_2}}$$

For two-dimensional data, this boundary is hence a straight line defined by:

$$P(y = 1|\mathbf{x}) = 0.5 \Leftrightarrow e^{-(\beta_0 + \beta_1 x_1 + \beta_2 x_2)} = 1 \Leftrightarrow 0 = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

This can also be rewritten as follows:

$$x_2 = -\frac{\beta_1}{\beta_2}x_1 - \frac{\beta_0}{\beta_2}$$

The coefficients determine both the orientation (slope) and position (intercept) of this separating line (the decision boundary).


# Decision Boundary Example

For example, let's assume we find the following parameters for the logistic regressor  $P(y = 1|x_1, x_2) = \sigma(\beta_0 + \beta_1 x_1 + \beta_2 x_2)$ :

$\beta_0$	$\beta_1$	$\beta_2$
-3.47	1.17	1.43

These parameters define the following decision boundary:

$$x_2 = -\frac{1.17}{1.43} + \frac{3.47}{1.43} = 2.43x_1 - 0.82$$



# The Softmax Regressor for Multiclass Problems

Softmax regression provides an alternative formulation to multinomial logistic regression that avoids selecting a baseline class. Instead of estimating K-1 coefficient sets, it estimates K sets symmetrically.

This formulation is widely used in machine learning and deep learning, where statistical interpretation of individual coefficients is less critical than predictive performance.

The softmax function ensures all class probabilities sum to 1:

$$P(Y = k | X = \mathbf{x}) = \frac{e^{\beta_k^T \mathbf{x}}}{\sum_{l=1}^K e^{\beta_l^T \mathbf{x}}}$$

Note that this provides an **overparametrization** of the model, as it includes more parameters than strictly needed.

- ❑ The Softmax regressor is optimized using a similar loss function to the Multinomial and Logistic regressors.

# Softmax Decision Boundaries

The core decision rule for a Softmax regressor is to assign each observation to the class with the highest predicted probability. Unlike multinomial logistic regression, which uses a baseline class, Softmax treats all classes symmetrically.


The **decision boundary** between any two classes  $i$  and  $j$  is formed where their probabilities are equal:

$$P(Y = i|\mathbf{x}) = P(Y = j|\mathbf{x})$$

Through mathematical simplification, this condition directly translates to the equality of their linear score functions:

$$\beta_i^T \mathbf{x} = \beta_j^T \mathbf{x} \Rightarrow (\beta_i - \beta_j)^T \mathbf{x} = 0$$

Geometrically, this equation defines a **hyperplane** in the feature space. In two dimensions, this is a straight line, and in three dimensions, it's a flat plane.



This means that Softmax regression effectively partitions the feature space into distinct regions, with each region corresponding to the class that is most likely according to the model. These partitions are always separated by linear boundaries.

## Linear Boundaries

Softmax regression generates **linear decision boundaries** between classes, similar to binary logistic regression, but extended to multiple classes.

## Score Equality

Each boundary is defined by the equality of two linear functions,  $\beta_i^T \mathbf{x}$  and  $\beta_j^T \mathbf{x}$ , for any pair of classes.

## Space Partitioning


The model systematically divides the feature space into regions, each assigned to the class with the highest probability score.

## Coefficient Role

The learned weights  $\beta_k$  for each class determine the precise orientation and position of these separating hyperplanes.

# From Binary to Multi-Class: One-vs-Rest (OvR)

## One-vs-all (one-vs-rest)



One-vs-Rest provides a technique to turn any binary classification model into a multi-class model. This is also known as "one vs all".

The OvR approach turns one K-class problem into K separate binary classification problems. To do this, we need a binary classifier (like Logistic Regression) that outputs a confidence score (a probability), which we can use to find the most likely class.

## Details

Given a classification task with K classes, the OvR approach works as follows:

1

### 1. Deconstruct the Problem

We train K independent binary classifiers, denoted as  $h_k(\mathbf{x})$ , for  $k \in \{1, \dots, K\}$ . Each classifier  $h_k$  is trained to distinguish class  $k$  from all other classes (the "rest").

Specifically, for each classifier  $h_k$ , the training set is relabeled:

- **Positive examples:** Instances belonging to class  $k$ .
- **Negative examples:** Instances belonging to any class  $j \neq k$ .

For instance:

- Classifier  $h_1$ : Distinguishes Class 1 from {Class 2, ..., Class  $K$ }
- Classifier  $h_2$ : Distinguishes Class 2 from {Class 1, Class 3, ..., Class  $K$ }
- ...
- Classifier  $h_K$ : Distinguishes Class  $K$  from {Class 1, ..., Class  $K - 1$ }

2

### 2. Classify a New Instance

To classify a new input instance  $\mathbf{x}$ , it is fed to all  $K$  trained binary classifiers. Each classifier  $h_k$  outputs a confidence score, typically a probability  $P(y = k|\mathbf{x})$ , indicating the likelihood that  $\mathbf{x}$  belongs to class  $k$ .

Mathematically, for each class  $k$ , we obtain:

$$s_k(\mathbf{x}) = P(y = k|\mathbf{x})$$

Where  $s_k(\mathbf{x})$  represents the confidence score for class  $k$ .

3



### 3. Select the Predicted Class

The final predicted class,  $\hat{y}$ , is determined by selecting the class whose corresponding classifier yielded the highest confidence score among all  $K$  classifiers.

$$\hat{y} = \operatorname{argmax}_{k \in \{1, \dots, K\}} s_k(\mathbf{x})$$

This decision rule assigns the instance  $\mathbf{x}$  to the class  $k$  for which its binary classifier  $h_k$  is most confident.

# Multiclass Logistic Regression in Python



# Conclusions and Next Steps



## We Have Explored:

- Multinomial logistic regression
- Geometrical interpretation of the coefficients of a logistic regressor
- Softmax regressor
- Geometrical interpretation of the coefficients of a softmax regressor
- One vs rest

In the next lectures, we will look at the Naive Bayes classifier

## References

- Chapter 4 of [1]
- <https://medium.com/@b.terryjack/tips-and-tricks-for-multi-class-classification-c184ae1c8ffc>

[1] James, Gareth Gareth Michael. An introduction to statistical learning: with applications in Python, 2023. <https://www.statlearning.com>